

Unmasking the Face: Emotion Detection using Machine Learning

by Maria Gil Rodriguez



Introduction



Dataset



Pre-processing



**Classic
Machine
Learning
Models**



Q&A



**Convolutional
Neural
Networks**



Future work



Conclusions

Introduction

- Human beings are responsible for the depiction of six basic emotions: happiness, anger, surprise, sadness, fear, and disgust.



- Creating a model that can accurately classify these emotions can be extremely useful in a variety of areas such as image processing, cybersecurity, robotics, psychological studies, virtual reality applications, etc.

Objectives

Objectives

The two main objective of this capstone are:

- Categorizing faces based on the emotion shown in the expression into one of seven categories (the six basic emotions plus one category for neutral)
- Studying the feature extraction with both supervised and unsupervised learning techniques

Dataset

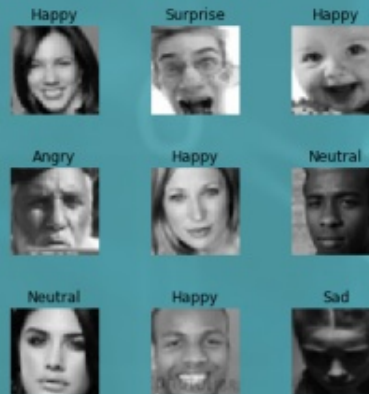
- FER2013 (Kaggle Competition)
- 48x48 pixel grayscale images of faces

- 3 sets:

Training	22968 images
Validation	5741 images
Test (no labels)	3589 images

- Classes:

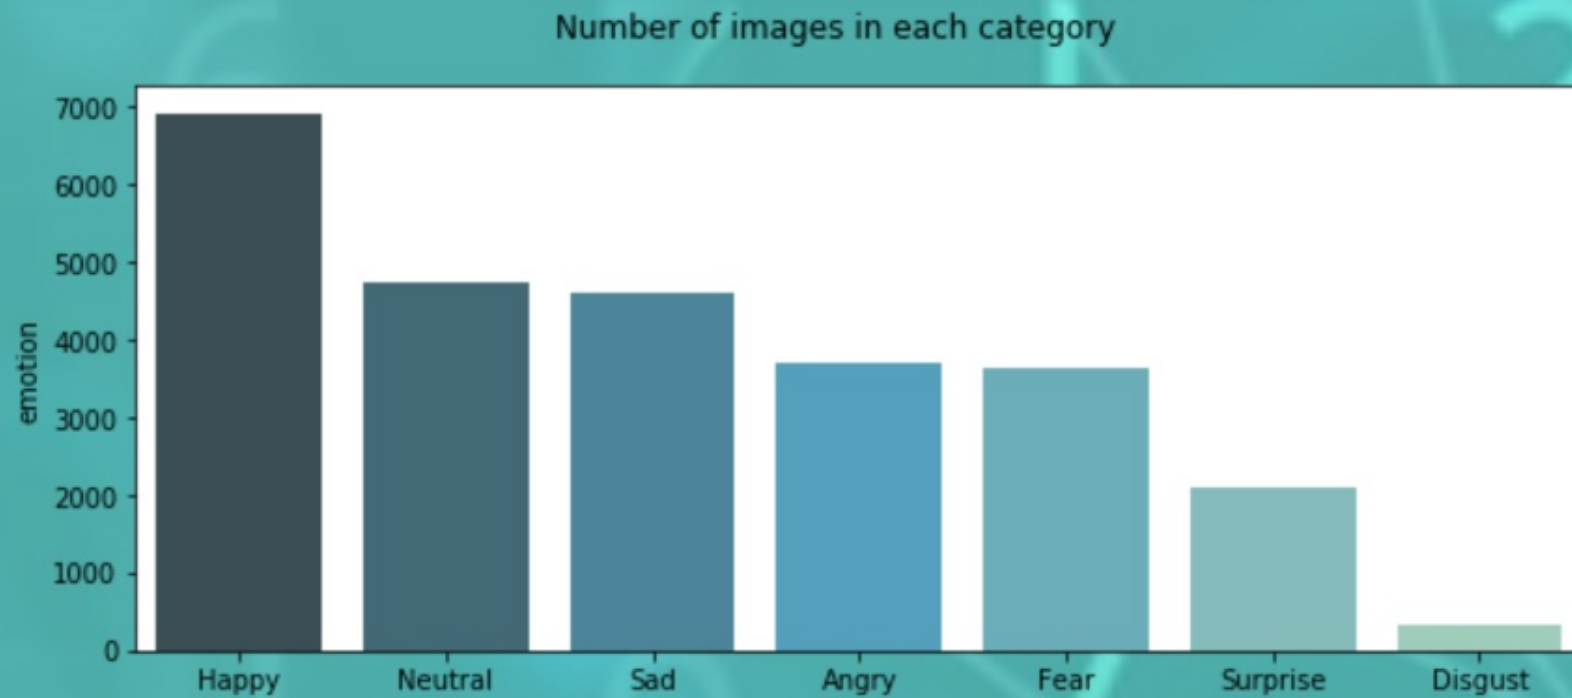
0=Angry
1=Disgust
2=Fear
3=Happy
4=Sad
5=Surprise
6=Neutral



***Class
imbalance***



Class imbalance



Pre-processing

	emotion	pixels
0	0	[70.0, 80.0, 82.0, 72.0, 58.0, 58.0, 60.0, 63....
1	0	[151.0, 150.0, 147.0, 155.0, 148.0, 133.0, 111...
2	2	[231.0, 212.0, 156.0, 164.0, 174.0, 138.0, 161...
3	4	[24.0, 32.0, 36.0, 30.0, 32.0, 23.0, 19.0, 20....
4	6	[4.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...

For the classic ML methods:

- Shuffle
- 100 images of each class for the training set
- 50 in the validation set

This:

- Avoids the class imbalance
- Reduces run time

Classic ML Methods

- Cross validation
- Grid Search

Random Forest	Gradient Boosting	SVC
0.25	0.26	0.25

PCA

LBP

PCA

300 components

Mean face for all the components



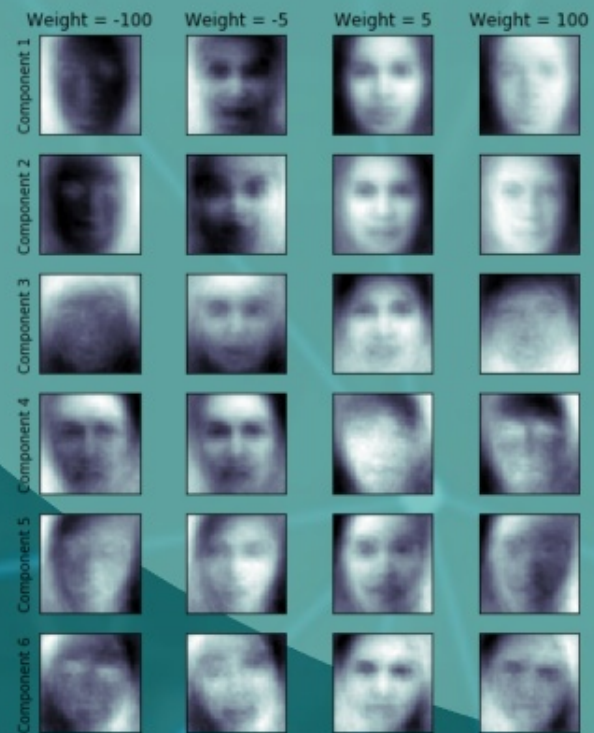
PCA

Eigen Faces

First 24 Eigen Faces



PCA



PCA

	Random Forest	Gradient Boosting	SVC
	0.25	0.26	0.25
PCA	0.23	0.23	0.26

Local Binary Patterns

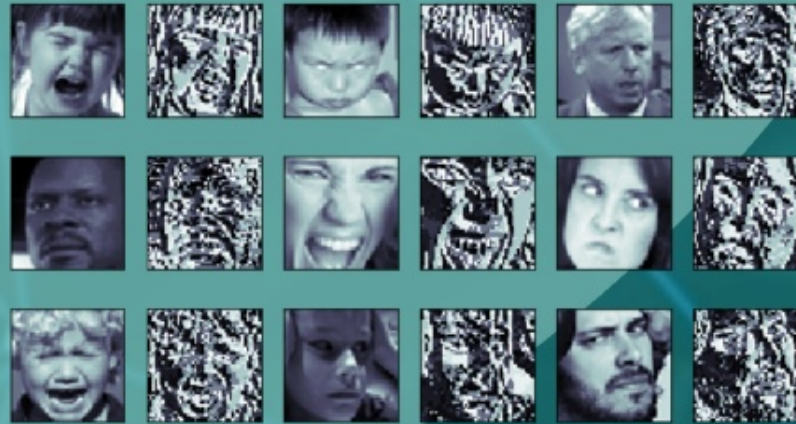
We consider the values of a 3x3 pixel neighborhood

- For each pixel in a cell, compare the pixel to each of its 8 neighbors (on its left-top, left-middle, left-bottom, right-top, etc.).
- Follow the pixels along a circle. Where the center pixel's value is greater than the neighbor's value, write "0". Otherwise, write "1".
- This gives an 8-digit binary number (which is usually converted to decimal for convenience).



Local Binary Patterns

Some examples of images before and after LBP



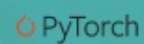
Local Binary Patterns

	Random Forest	Gradient Boosting	SVC
	0.25	0.26	0.25
PCA	0.23	0.23	0.26
LBT	0.18	0.19	0.21

CNNs

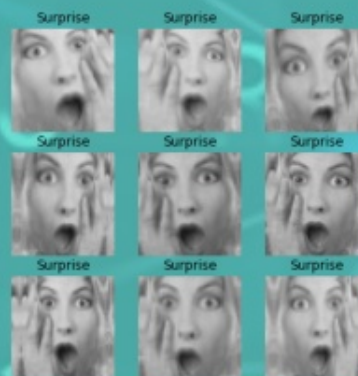


fast.ai library
PyTorch



Data Augmentation:

- Rotation
- Zoom
- Lighting
- Warp
- Padding
- Horizontal flip



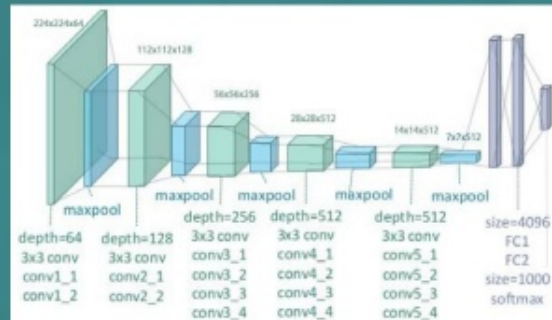
***Pre-trained
models***

My model

***Saving the
model***

Pre- trained Models

VGG-19: Transfer learning

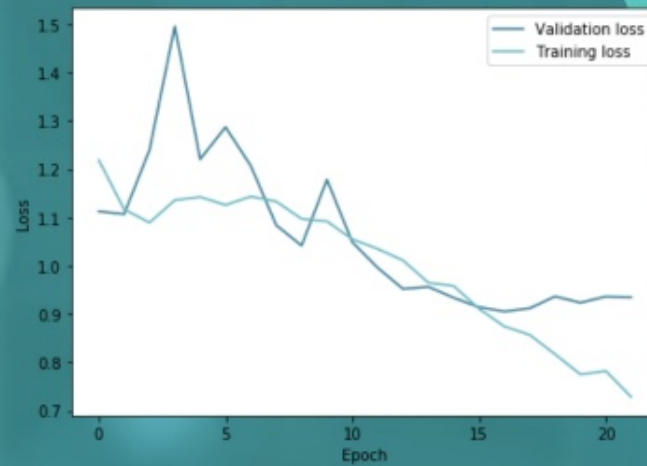


Tuning:

- # Epochs
- Learning rate

Pre- trained Models

VGG-19: Transfer learning



Pre- trained Models

VGG-19: Transfer learning

Confusion matrix

Actual	Angry	478	8	96	35	92	95	6
	Disgust	10	50	4	3	1	10	2
	Fear	73	7	436	21	89	163	66
	Happy	20	1	32	1267	72	27	20
	Neutral	38	1	38	77	706	127	16
	Sad	68	4	93	33	163	572	6
	Surprise	16	1	72	38	15	6	467
		Predicted						

Accuracy: 69%

Most confused:

Fear / Sad: 163

Sad / Neutral: 163

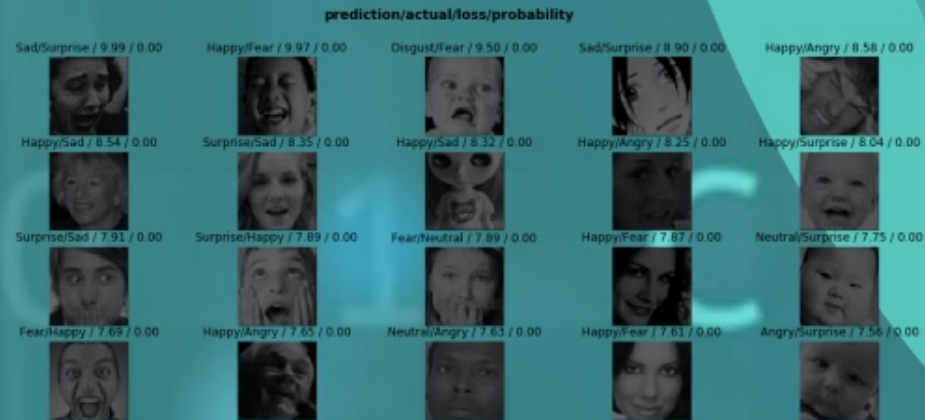
Neutral / Sad: 127

Angry / Fear: 96

Angry / Sad: 95

Pre- trained Models

VGG-19: Transfer learning

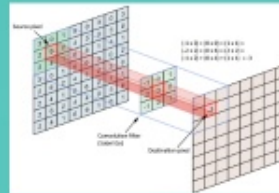


Cleaning?

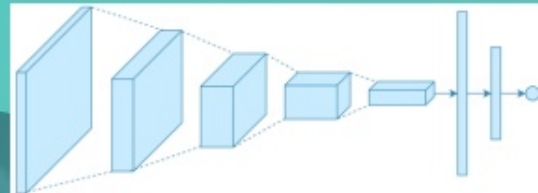
My model

Model from scratch

- Convolutions (stride and padding)

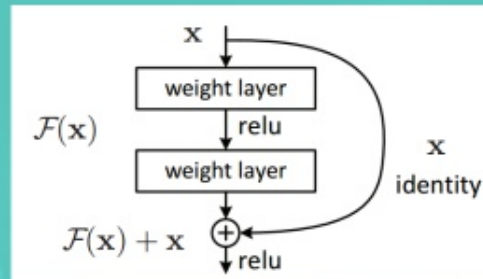


- RELU
- Pooling
- Batch Normalization



My model

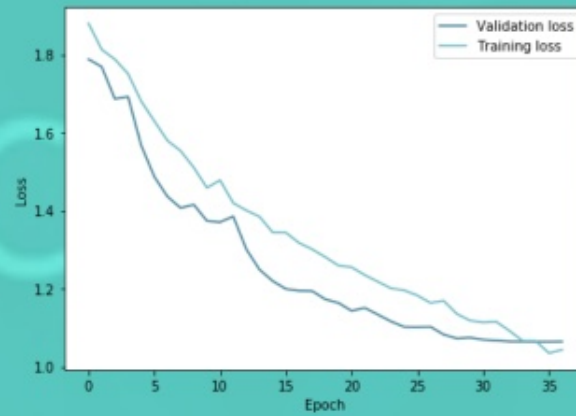
- Residual blocks



```
model_own = nn.Sequential(  
    conv_and_res(3, 8), #24  
    conv_and_res(8, 16), #12  
    conv_and_res(16, 32), #6  
    conv_and_res(32, 64), #3  
    conv2(64, 128), #2  
    nn.MaxPool2d(2, 2),  
    conv2(128, 256), #1  
    Flatten(),  
    nn.Linear(256, 128),  
    nn.Dropout(.5),  
    nn.Linear(128, 7))
```

- 14 convolutional layers
- 2 fully connected layers

My model



61% accuracy

My model

- CNN Layers Visualization

Surprise



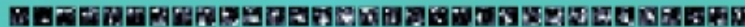
FeatureBlock number 1



FeatureBlock number 2



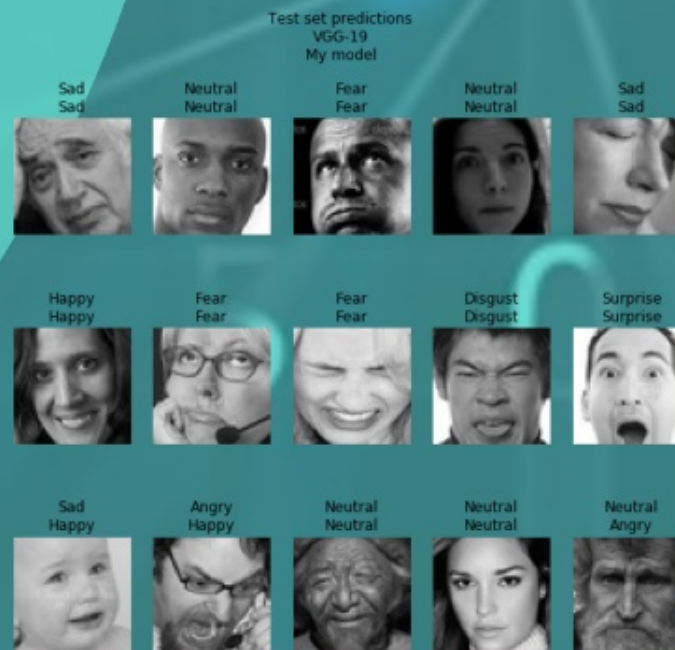
FeatureBlock number 3



FeatureBlock number 4



Saving the model



Conclusions

- We were able to successfully extract the image feature with unsupervised learning techniques. However, the classic supervised learning classification models performed poorly.
- Pre-trained CNN models performed pretty well and fast (low number of epochs). We were able to classify emotions with 69% accuracy.
- We were able to create a CNN from scratch that performed pretty well, with an accuracy of . However, the training process was longer, needing more epochs to achieve the aforementioned accuracy.

Future Work

- Use other datasets with better labels and higher image resolution.
- In the pre-trained models, we can introduce weights obtained training facial features of other datasets (e.g. VGG-Face).
- Improve our own model: changing achitecture and tuning hyperparameters.
- Use the model to predict facial emotions in videos.

Q & A

